Lecture Notes on Operating Systems

Lab: Building a Shell

Goal

The goal of this lab is to develop a simple shell to execute user commands, much like the `bash` shell in Linux. This lab will deepen your understanding of various concepts of process management in Linux.

Before you begin

- Familiarize yourself with the various process related system calls in Linux: `fork`, `exec`, `exit` and `wait`. Understand the different variants of these system calls. The "man pages" in Linux are a good source of learning. You can access the man pages from the Linux terminal by typing `man fork`, `man 2 fork` and so on. You can also find several helpful links online (e.g., `http://manpages.ubuntu.com/manpages/trusty/man2/fork.2.html`).

- Familiarize yourself with simple built-in commands in Linux like `echo`, `cat`, `sleep`, `ls`, `ps`, `top`, `grep` and so on.

- Understand the `chdir` system call in Linux (see `man chdir`). This will be useful to implement the `cd` command in your shell.

- Understand the concepts of foreground and background execution in Linux. Execute various commands on the Linux shell to understand the behavior of these modes of execution.

- Understand signals and signal handling in Linux. Understand how processes can send signals to one another using the `kill` system call. Read up on how to write custom signal handlers to "catch" signals and override the default signal handling mechanism, using interfaces such as `signal()` or `sigaction()`.

- Understand the notion of processes and process groups. When a signal like Ctrl+C is sent to a process, it is delivered to all processes in its process group, including all its children. If you want only some subset of the children of a process (and not the others) to receive a signal, you must put these children in a separate process group, say, by using the `setpgid` system call. Lookup this system call in the man pages to learn more about it.

Part A: A simple shell

We will first build a simple shell to mostly run Linux built-in commands. A shell takes in user input, forks one or more child processes using the `fork` system call, calls `exec` from these children to execute user commands, and reaps the dead children using the `wait` system call. Your shell must execute all simple Linux commands like `ls, cat, echo` and `sleep`. These commands are readily available as executables on Linux, and your shell must simply invoke the existing executable.

Your simple shell must use the string "$ " as the command prompt. Your shell must run in one of two modes: interactive or batch. If a batch file of commands is provided as command-line input to your program, then your shell must execute all commands in the batch file one after the other. You are provided a sample batch file `commands.txt` as an example. If no command-line argument is provided, your shell should interactively accept inputs from the user and execute them.

In this part, the shell must return for user input (or move on to the next command in the batch) only after the execution of the previous command completes. Further, in this part, a shell in interactive mode should continue execution indefinitely until the user hits Ctrl+C to terminate the shell. In batch mode, the shell must exit once it reaches the end of the batch file.

You can assume that the command to run and its arguments are separated by one or more spaces in the input, so that you can "tokenize" the input stream using spaces as the delimiters. For this part, you can assume that the Linux built-in commands are invoked with simple command-line arguments, and without any special modes of execution like background execution, I/O redirection, or pipes. You need not parse any other special characters in the input stream.

A skeleton code `my_shell.c` is provided to get you started. This program reads input (interactively or from a batch) and tokenizes the input for you. You must add code to this file to execute the commands found in the "tokens". You may assume that the input command has no more than 1024 characters, and no more than 64 tokens. Further, you may assume that each token is no longer than 64 characters. You can compile and run this code in two ways as shown below.

- `./my_shell` will run the program in interactive mode.

- `./my_shell commands.txt` will run the program in batch mode.

Once you complete the execution of the built-in commands, proceed to implement support for the simple `cd` command in your shell. The command `cd <directoryname>` must cause the shell process to change its working directory, and `cd ..` should take you to the parent directory. You need not support other variants of `cd` that are available in the various Linux shells. For example, just typing `cd` will take you to your home directory in some shells; you need not support such complex features.

Note that for all commands you implement in this lab, an incorrect number of arguments or incorrect command format should print an error message **Shell: Incorrect command** to the display. Further, any error messages displayed during the execution of the commands should also be printed to the display. After such errors are printed by the shell, the shell should not crash. It must simply move on and prompt the user for the next command.

Part B: Serial, parallel, and background execution

Now, we will build support for executing multiple commands at a time in your shell, as described below.

- If a command is followed by &, the command must be executed in the background. That is, the shell must start the execution of the command, and return to prompt the user for the next input, without waiting for the previous command to complete. The output of the command can get printed to the shell as and when it appears.

- Multiple user commands separated by && should be executed one after the other in sequence in the foreground. The shell must move on to the next command in the sequence only after the previous one has completed (successfully, or with errors). An error in one command should cause the shell to print the error message **Shell: Incorrect command** and move on to the next command. The shell should return to the command prompt after all the commands in the sequence have finished execution.

- Multiple commands separated by &&& should be executed in parallel in the foreground. That is, the shell should start execution of all commands simultaneously, and return to command prompt after all commands have finished execution.

Across all cases, carefully ensure that the shell reaps all its children that have terminated. For commands that must run in the foreground, the shell must wait for and reap its terminated foreground child processes before it prompts the user for the next input. For the command that creates background child processes, the shell must periodically check and reap any terminated background processes while running other commands. When the shell reaps a terminated background process at a future time, it must print a message **Shell: Background process finished** to let the user know that a background process has finished.

In all the cases above, you may assume that each of the individual commands are simple Linux built-in commands without pipes or redirections or any other special cases. You may also assume that there are spaces on either side of the special tokens like &, &&, and &&&. You may assume that there are no more than 64 foreground or background commands executing at any given time.

Part C: Signal handling and exit

Up until now, your shell executes in an infinite loop, and only the signal SIGINT (Ctrl+C) would have caused it to terminate. Now, you will modify your shell so that the signal SIGINT does not terminate it. You will also implement the `exit` command that will cause the shell to terminate its infinite loop and exit.

When the user hits Ctrl+C, the shell must correctly handle the signal and not terminate itself. Instead, it should terminate the current foreground processes (the current command in serial execution, or all the commands in a parallel execution), and return to the command prompt. The background processes should remain unaffected by the SIGINT. When the shell receives the `exit` command, it must terminate all background processes, clean up any internal state (e.g., free dynamically allocated memory), and finally terminate.

Recall that, by default, any signal like SIGINT will be delivered to the shell and all its children. To solve this part correctly, you must carefully place some children in a different process group and send signals only to those processes that must receive it.

Submission instructions

- You must submit the shell code `my_shell.c` or `my_shell.cpp`.

- Place this file and any other files you wish to submit in your submission directory, with the directory name being your roll number (say, 12345678).

- Tar and gzip the directory using the command `tar -zcvf 12345678.tar.gz 12345678` to produce a single compressed file of your submission directory. Submit this tar gzipped file on Moodle.

NOTE: Please follow all instructions carefully, especially those pertaining to the format of the various messages (e.g., error messages) that must be printed to the screen.